

Collision Avoidance (ResNet18)

Reference: [NVIDIA JetBot Github Repository](#)

0. Import modules

```
In [1]: import torch
import torch.optim as optim # for SGD
import torch.nn.functional as F
from torch.utils.data import random_split, DataLoader

import torchvision
import torchvision.models as models # for resnet18
import torchvision.transforms as transforms
from torchvision.datasets import ImageFolder

import os # isdir, mkdir
import matplotlib.pyplot as plt
from time import strftime, localtime
```

1. Prepare the dataset

```
In [2]: # DATASET_PATH = "./datasets/dataset_white"
DATASET_PATH = "./datasets/dataset_blue"

IMAGE_WIDTH = 224
IMAGE_HEIGHT = 224

# The two constants below use the values specified in the reference.
NORMALIZE_MEAN = (0.485, 0.456, 0.406)
NORMALIZE_STD = (0.229, 0.224, 0.225)

total_dataset = ImageFolder(
    DATASET_PATH,
    transforms.Compose([
        transforms.ColorJitter(0.1, 0.1, 0.1, 0.1),
        transforms.Resize((IMAGE_HEIGHT, IMAGE_WIDTH)),
        transforms.ToTensor(),
        transforms.Normalize(NORMALIZE_MEAN, NORMALIZE_STD)
    ])
)

print(f"{len(total_dataset)} images have been loaded.") # Logger
```

240 images have been loaded.

```
In [3]: SPLIT_RATIO = (0.8, 0.1, 0.1) # train : valid : test

total_data_num = len(total_dataset)

train_data_num = int(total_data_num * SPLIT_RATIO[0])
valid_data_num = int(total_data_num * SPLIT_RATIO[1])
model_data_num = train_data_num + valid_data_num

test_data_num = int(total_data_num * SPLIT_RATIO[2])

model_dataset, test_dataset = random_split(total_dataset, (model_data_num, test_data_num))
```

```

train_dataset, valid_dataset = random_split(model_dataset, (train_data_num, v

#-- Logger --#
print(f"Train Dataset: {len(train_dataset)} images.") # print(train_data_num,
print(f"Validation Dataset: {len(valid_dataset)} images.") # print(valid_data
print(f"Test Dataset: {len(test_dataset)} images.") # print(test_data_num)
#-- Logger --#

```

Train Dataset: 192 images.
 Validation Dataset: 24 images.
 Test Dataset: 24 images.

In [4]:

```

BATCH_SIZE = 8

train_loader = DataLoader(
    train_dataset,
    batch_size = BATCH_SIZE,
    shuffle = True,
    num_workers = 0
)

valid_loader = DataLoader(
    train_dataset,
    batch_size = BATCH_SIZE,
    shuffle = True,
    num_workers = 0
)

```

2. Define the model (ResNet18)

In [5]:

```

model = models.resnet18(pretrained=True)
model.fc = torch.nn.Linear(512, 2) # 2 for 'blocked' and 'free'

device = torch.device('cpu')
if torch.cuda.is_available():
    device = torch.device('cuda')
    print("This environment supports the CUDA.") # Logger
else:
    print("This environment does not support the CUDA.") # Logger
    print("The model will be running on the CPU instead.") # Logger
    # pass

model = model.to(device)

# print(model)

```

This environment supports the CUDA.

3. Train the model

In [6]:

```

if not os.path.isdir("./best_models"):
    os.mkdir("./best_models")

CURRENT_TIME = strftime('%Y%m%d_%H%M%S', localtime())
BEST_MODEL_PATH = f"./best_models/best_model_resnet18_{CURRENT_TIME}.pth"

# hyper parameters
EPOCHS = 30
LEARNING_RATE = 0.001
MOMENTUM = 0.9

```

```

L2_CONST = 1e-4

best_accuracy = 0.0 # validation accuracy

# SGD optimizer with L2 regularization
optimizer = optim.SGD(model.parameters(),
                        lr=LEARNING_RATE,
                        momentum=MOMENTUM,
                        weight_decay=L2_CONST)

accuracy_history = []

EPOCH_DIGIT = len(str(EPOCHS)) # for Logger

```

```

In [7]: # model training loop
for epoch in range(EPOCHS):

    for images, labels in iter(train_loader):
        images = images.to(device)
        labels = labels.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = F.cross_entropy(outputs, labels)
        loss.backward()
        optimizer.step()

    valid_error = 0.0
    for images, labels in iter(valid_loader):
        images = images.to(device)
        labels = labels.to(device)
        outputs = model(images)
        outputs = outputs.argmax(1)
        valid_error += float(torch.sum(torch.abs(labels - outputs)))

    valid_accuracy = 1.0 - (valid_error / valid_data_num)

    if valid_accuracy < 0:
        valid_accuracy = 0

    accuracy_history.append(valid_accuracy)

    print(f"[Epoch {epoch : >{EPOCH_DIGIT}d}] Validation accuracy: {valid_acc

    if valid_accuracy > best_accuracy:
        print("\tSave the best model") # Logger
        torch.save(model.state_dict(), BEST_MODEL_PATH)
        best_accuracy = valid_accuracy

print("Training Complete!") # Logger
print(f"Best validation accuracy: {best_accuracy: .5f}") # Logger

```

```

[Epoch 0] Validation accuracy: 0.87500
Save the best model
[Epoch 1] Validation accuracy: 0.95833
Save the best model
[Epoch 2] Validation accuracy: 0.83333
[Epoch 3] Validation accuracy: 0.79167
[Epoch 4] Validation accuracy: 0.91667
[Epoch 5] Validation accuracy: 0.95833
[Epoch 6] Validation accuracy: 1.00000
Save the best model
[Epoch 7] Validation accuracy: 0.83333
[Epoch 8] Validation accuracy: 0.91667
[Epoch 9] Validation accuracy: 1.00000

```

```

[Epoch 10] Validation accuracy: 1.00000
[Epoch 11] Validation accuracy: 0.95833
[Epoch 12] Validation accuracy: 0.95833
[Epoch 13] Validation accuracy: 1.00000
[Epoch 14] Validation accuracy: 1.00000
[Epoch 15] Validation accuracy: 1.00000
[Epoch 16] Validation accuracy: 1.00000
[Epoch 17] Validation accuracy: 0.91667
[Epoch 18] Validation accuracy: 0.87500
[Epoch 19] Validation accuracy: 0.91667
[Epoch 20] Validation accuracy: 0.95833
[Epoch 21] Validation accuracy: 1.00000
[Epoch 22] Validation accuracy: 1.00000
[Epoch 23] Validation accuracy: 1.00000
[Epoch 24] Validation accuracy: 1.00000
[Epoch 25] Validation accuracy: 1.00000
[Epoch 26] Validation accuracy: 1.00000
[Epoch 27] Validation accuracy: 1.00000
[Epoch 28] Validation accuracy: 1.00000
[Epoch 29] Validation accuracy: 0.83333
Training Complete!
Best validation accuracy: 1.00000

```

```

In [8]: if not os.path.isdir("./plots"):
        os.mkdir("./plots")

        PLOT_PATH = f"./plots/validation_accuracy_plot_resnet18_{CURRENT_TIME}.png"


        title = "Validation Accuracy Plot (ResNet18)"
        subtitle = f"(learning_rate={LEARNING_RATE}, momentum={MOMENTUM}, l2_constant={L2_CONSTANT})"

        plt.plot(accuracy_history)

        plt.suptitle(title)
        plt.title(subtitle)
        plt.xlabel("Epoch")
        plt.ylabel("Validation accuracy")

        plt.savefig(PLOT_PATH)
        plt.show()

```



4. Test the model

```

In [9]: model = models.resnet18(pretrained=False)
        model.fc = torch.nn.Linear(512, 2)

```

```

print(f"Load model from \"{BEST_MODEL_PATH}\".") # Logger

model.load_state_dict(torch.load(BEST_MODEL_PATH))

model = model.to(device)
model = model.eval().half()

```

Load model from "./best_models/best_model_resnet18_20210204_221516.pth".

In [10]:

```

test_loader = DataLoader(
    test_dataset,
    batch_size=1,
    shuffle=True,
    num_workers=0
)

```

In [11]:

```

correct_case_count = 0

for case, sample in enumerate(iter(test_loader)):
    image, label = sample
    image = image.to(device).half()
    label = int(label)
    predict = model(image)
    predict = F.softmax(predict, dim=1)
    predict = predict.flatten()

    #-- Logger --#
    print(f"[Test Case {case}]")
    print(f"\t[Prediction] {float(predict[0]): .5f} : {float(predict[1]): .5f}")
    # print(f"\t[Prediction] Blocked : Free")
    print(f"\t[Real output] {label}") # 0: Blocked, 1: Free
    #-- Logger --#

    if label == 1 and float(predict[0]) < float(predict[1]):
        correct_case_count += 1
        print(f"\t[Result] Correct") # Logger
    elif label == 0 and float(predict[0]) > float(predict[1]):
        correct_case_count += 1
        print(f"\t[Result] Correct") # Logger
    else:
        print(f"\t[Result] Incorrect") # Logger
        # pass

print(f"[Total Test Accuracy] {correct_case_count/test_data_num : .5f}")

```

```

[Test Case 0]
    [Prediction] 0.99756 : 0.00250
    [Real output] 0
    [Result] Correct
[Test Case 1]
    [Prediction] 0.00004 : 1.00000
    [Real output] 1
    [Result] Correct
[Test Case 2]
    [Prediction] 0.00037 : 0.99951
    [Real output] 1
    [Result] Correct
[Test Case 3]
    [Prediction] 1.00000 : 0.00000
    [Real output] 0
    [Result] Correct
[Test Case 4]
    [Prediction] 0.99756 : 0.00252

```

```
[Real output] 0
[Result] Correct
[Test Case 5]
[Prediction] 0.00010 : 1.00000
[Real output] 1
[Result] Correct
[Test Case 6]
[Prediction] 0.00001 : 1.00000
[Real output] 1
[Result] Correct
[Test Case 7]
[Prediction] 0.00002 : 1.00000
[Real output] 1
[Result] Correct
[Test Case 8]
[Prediction] 0.00004 : 1.00000
[Real output] 1
[Result] Correct
[Test Case 9]
[Prediction] 1.00000 : 0.00009
[Real output] 0
[Result] Correct
[Test Case 10]
[Prediction] 0.00070 : 0.99951
[Real output] 1
[Result] Correct
[Test Case 11]
[Prediction] 0.00029 : 0.99951
[Real output] 1
[Result] Correct
[Test Case 12]
[Prediction] 1.00000 : 0.00013
[Real output] 0
[Result] Correct
[Test Case 13]
[Prediction] 0.00002 : 1.00000
[Real output] 1
[Result] Correct
[Test Case 14]
[Prediction] 0.99854 : 0.00153
[Real output] 0
[Result] Correct
[Test Case 15]
[Prediction] 0.99316 : 0.00674
[Real output] 0
[Result] Correct
[Test Case 16]
[Prediction] 0.00002 : 1.00000
[Real output] 1
[Result] Correct
[Test Case 17]
[Prediction] 0.00015 : 1.00000
[Real output] 1
[Result] Correct
[Test Case 18]
[Prediction] 1.00000 : 0.00018
[Real output] 0
[Result] Correct
[Test Case 19]
[Prediction] 0.00158 : 0.99854
[Real output] 1
[Result] Correct
[Test Case 20]
[Prediction] 0.00109 : 0.99902
[Real output] 1
[Result] Correct
[Test Case 21]
[Prediction] 1.00000 : 0.00013
[Real output] 0
```

```
      [Result] Correct
[Test Case 22]
      [Prediction] 0.97656 : 0.02328
      [Real output] 0
      [Result] Correct
[Test Case 23]
      [Prediction] 0.99951 : 0.00042
      [Real output] 0
      [Result] Correct
[Total Test Accuracy] 1.00000
```